


Dr. Seemab latif

Lecture 3


1 Oct 2024

AI: Representation and Problem Solving

Informed Search



Uninformed vs Informed Search



Today




- Informed Search
 - Heuristics
 - Greedy Search
 - A* Search



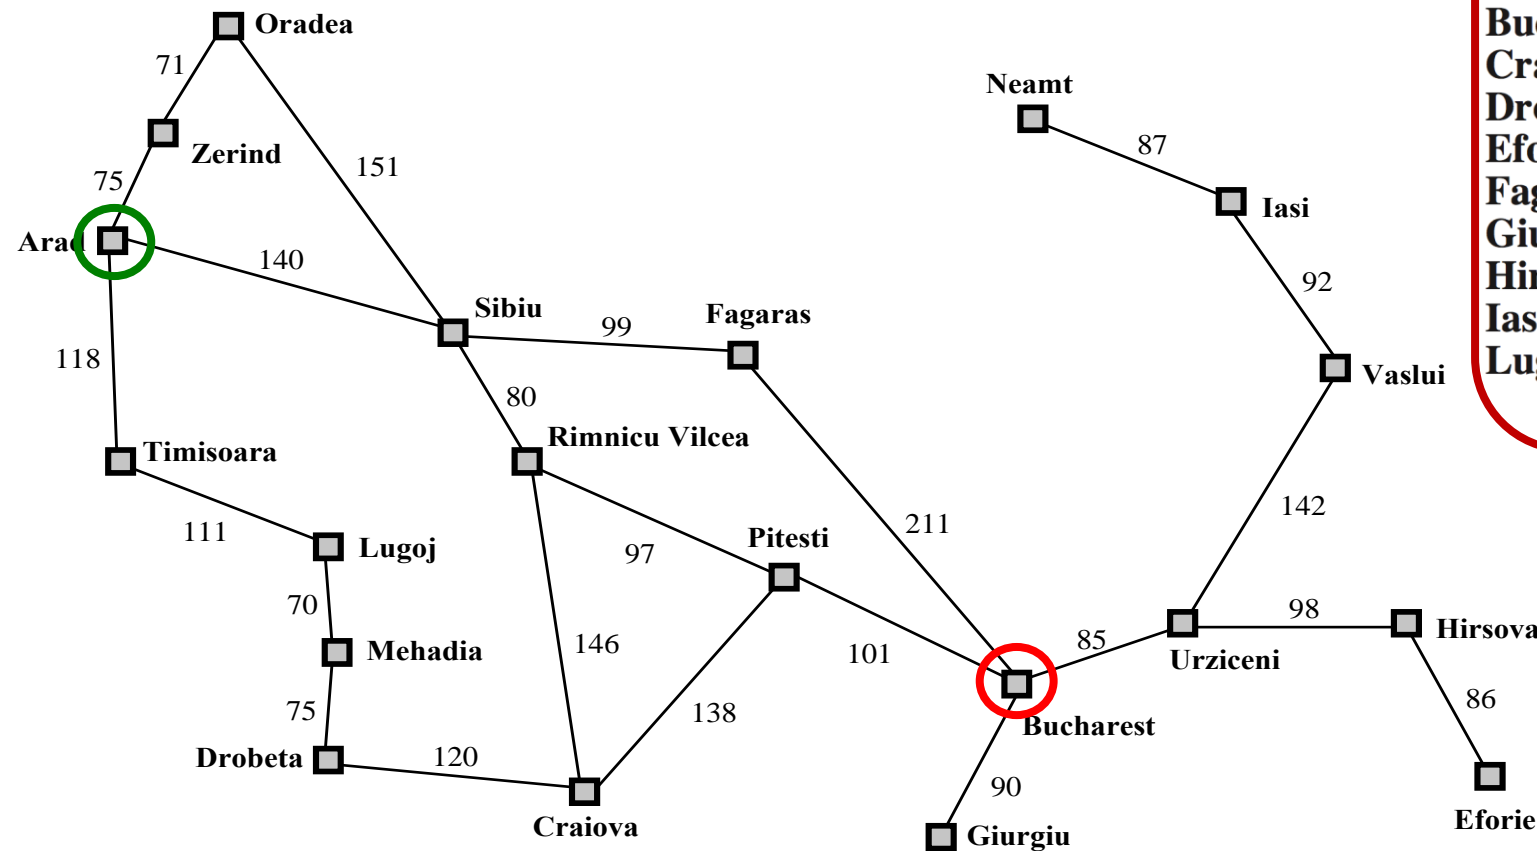
Search Heuristics

A heuristic is:

- A function that *estimates* how close a state is to a goal
- Designed for a particular search problem
- Examples: Manhattan distance, Euclidean distance for pathing



Example: Euclidean distance to Bucharest




Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374


$h(\text{state}) \rightarrow \text{value}$

Effect of heuristics

Guide search *towards the goal* instead of *all over the place*



Informed




Uninformed

Greedy Search




Greedy Search

- Expand the node that seems closest...(order frontier by h)
- What can possibly go wrong?




Sibiu-Fagaras-Bucharest =
 $99 + 211 = 310$

Sibiu-Rimnicu Vilcea-Pitesti-Bucharest =
 $80 + 97 + 101 = 278$



Greedy Search

- Strategy: expand a node that *seems* closest to a goal state, according to h
- Problem 1: it chooses a node even if it's at the end of a very long and winding road
- Problem 2: it takes h literally even if it's completely wrong



A* Search



A* Search



UCS



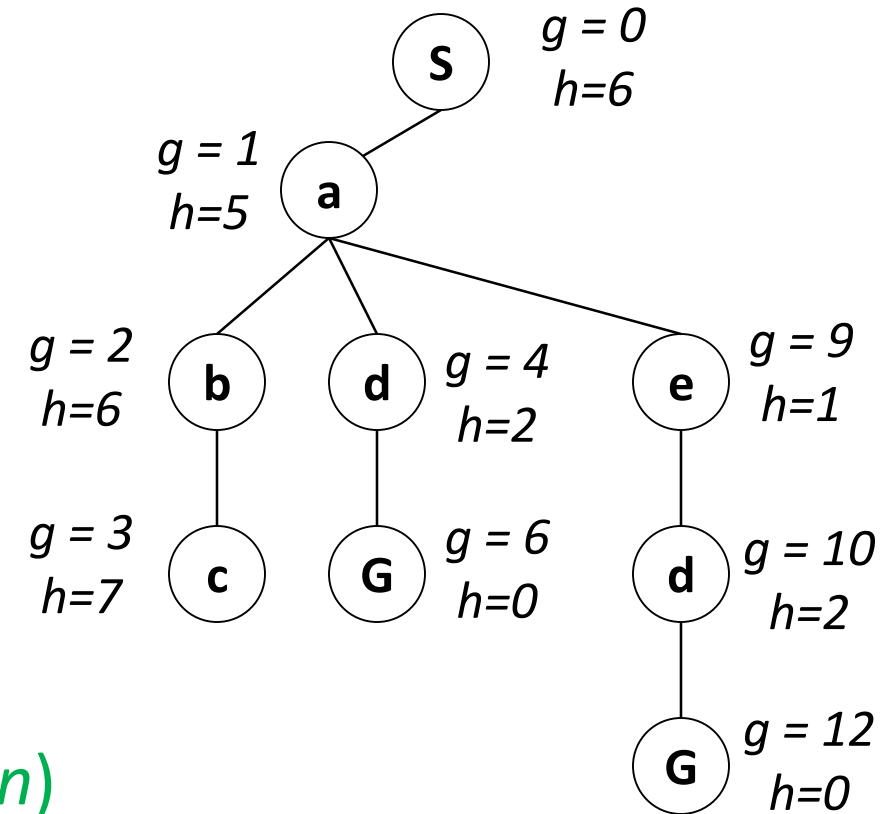

Greedy



A*


Combining UCS and Greedy

- **Uniform-cost** orders by path cost, or *backward cost* $g(n)$
- **Greedy** orders by goal proximity, or *forward cost* $h(n)$



- **A* Search** orders by the sum: $f(n) = g(n) + h(n)$

A* Search orders by the sum: $f(n) = g(n) + h(n)$




Current	Open list (Greedy)	Close
S(6)	A(5)	-
A(5)	E(1), D(2), B(6)	S(6)
E(1)	D(2), B(6)	S(6), A(5)
D(2)	G(0), B(6)	S(6), A(5), E(1)
G(0)	Goal	S(6), A(5), E(1), D(2)
Path =	S-A-E-D-G	
Cost =	1+8+1+2 = 12	

Current	Open list (UCS)	Close
S(0)	A(1)	
A(1)	B(2), D(4), E(9)	S(0)
B(2)	C(3) , D(4), E(9)	S(0), A(1)
C(3)	D(4), E(9)	S(0), A(1), B(2)
D(4)	G(6), E(9)	S(0), A(1), B(2), C(3)
G(6) = goal node	E(9)	S(0), A(1), B(2), C(3) D(4)

Current	Open list (A*)	Close
S(0+6=6)	A(1+5=6)	
A(6)	D(6), B(8), E(10)	S(6)
D(6)	G(6), B(8), E(10)	S(6), A(6)
G(6) = goal node	B(8), E(10)	S(6), A(6), D(6)
Path =	S-A-D-G	
Cost = 6		

A^*



function UNIFORM-COST-SEARCH(**problem**) **returns** a solution, or failure

initialize the **explored set** to be empty

initialize the **frontier** as a priority queue using $g(n)$ as the priority

add initial state of **problem** to **frontier** with priority $g(S) = 0$

loop do

if the **frontier** is empty **then**

return failure

 choose a **node** and remove it from the **frontier**

if the **node** contains a goal state **then**

return the corresponding solution

 add the **node** state to the **explored set**

 for each resulting **child** from node

if the **child** state is not already in the **frontier** or **explored set** **then**

 add **child** to the **frontier**

else if the **child** is already in the **frontier** with higher $g(n)$ **then**

 replace that **frontier** node with **child**

function A-STAR-SEARCH(**problem**) **returns** a solution, or failure

initialize the **explored set** to be empty

initialize the **frontier** as a priority queue using $f(n) = g(n) + h(n)$ as the priority

add initial state of **problem** to **frontier** with priority $f(S) = 0 + h(S)$

loop do

if the **frontier** is empty **then**

return failure

 choose a **node** and remove it from the **frontier**

if the **node** contains a goal state **then**

return the corresponding solution

 add the **node** state to the **explored set**

 for each resulting **child** from node

if the **child** state is not already in the **frontier** or **explored set** **then**

 add **child** to the **frontier**

else if the **child** is already in the **frontier** with higher $f(n)$ **then**

 replace that **frontier** node with **child**

A* Search Algorithms

- A* Tree Search

- Same tree search algorithm but with a **frontier** that is a priority queue using priority $f(n) = g(n) + h(n)$



- A* Graph Search

- Same as **UCS** graph search algorithm but with a **frontier** that is a priority queue using priority $f(n) = g(n) + h(n)$

UCS vs A* Contours

UCS vs A* Contours

A* expands mainly toward the goal, but does hedge its bets to ensure optimality



Comparison



Greedy




Uniform Cost




A*

Is A* Optimal?



- What went wrong?
- **Actual** bad goal cost < **estimated** good goal cost
- We need estimates to be less than actual costs!

Admissible Heuristics




Admissible Heuristics

- A heuristic h is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$


where $h^*(n)$ is the true cost to a nearest goal

- Example:




- Coming up with admissible heuristics is most of what's involved in using A^* in practice.

Optimality of A* Tree Search




A* Tree Search

State space graph




Search tree



Optimality of A* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible




Claim:

- A will be chosen for exploration (popped off the frontier) before B

Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the frontier
- Some ancestor n of A is on the frontier, too (Maybe the start state; maybe A itself!)
- Claim: n will be explored before B
 1. $f(n)$ is less than or equal to $f(A)$



$$f(n) = g(n) + h(n)$$

$$f(n) \leq g(A)$$

$$g(A) = f(A)$$

Definition of f -cost


Admissibility of h

$h = 0$ at a goal

Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the frontier
- Some ancestor n of A is on the frontier, too (Maybe the start state; maybe A itself!)
- Claim: n will be explored before B
 1. $f(n)$ is less than or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$




$$g(A) < g(B)$$
$$f(A) < f(B)$$

Suboptimality of B
 $h = 0$ at a goal

Optimality of A* Tree Search: Blocking


Proof:

- Imagine B is on the frontier
- Some ancestor n of A is on the frontier, too (Maybe the start state; maybe A itself!)
- Claim: n will be explored before B
 1. $f(n)$ is less than or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$
 3. n is explored before B
- All ancestors of A are explored before B
- A is explored before B
- A* search is optimal




$$f(n) \leq f(A) < f(B)$$

Optimality of A* Graph Search





What paths does A* graph search consider during its search?




- A) ~~S~~, ~~S-A~~, ~~S-C~~, S-C-G
- B) ~~S~~, ~~S-A~~, S-C, ~~S-A-C~~, S-C-G
- C) ~~S~~, ~~S-A~~, ~~S-A-C~~, S-A-C-G
- D) ~~S~~, ~~S-A~~, ~~S-C~~, ~~S-A-C~~, S-A-C-G



1-answer

What paths does A* graph search consider during its search?



A) ~~S~~, ~~S-A~~, ~~S-C~~, S-C-G





B) ~~S~~, ~~S-A~~, ~~S-C~~, ~~S-A-C~~, S-C-G

C) ~~S~~, ~~S-A~~, ~~S-A-C~~, S-A-C-G

D) ~~S~~, ~~S-A~~, ~~S-C~~, ~~S-A-C~~, S-A-C-G


A* Graph Search

What does the resulting graph tree look like?




A* Graph Search Gone Wrong?

State space graph




Search tree



Simple check against explored set blocks C


Fancy check allows new C if cheaper than old
but requires recalculating C's descendants

Admissibility of Heuristics



- Main idea: Estimated heuristic values \leq actual costs
 - Admissibility:
heuristic value \leq actual cost to goal
 $h(A) \leq$ actual cost from A to G


Consistency of Heuristics



- Main idea: Estimated heuristic costs \leq actual costs
 - Admissibility:
heuristic cost \leq actual cost to goal
 $h(A) \leq \text{actual cost from A to G}$
 - Consistency:
“heuristic step cost” \leq actual cost for each step
 $h(A) - h(C) \leq \text{cost}(A \text{ to } C)$
triangle inequality
 $h(A) \leq \text{cost}(A \text{ to } C) + h(C)$
- Consequences of consistency:
 - The f value along a path never decreases
 - A* graph search is optimal


Optimality of A* Graph Search

- Sketch: consider what A* does with a consistent heuristic:
 - Fact 1: In tree search, A* expands nodes in increasing total **f** value (**f-contours**)
 - Fact 2: For every state **s**, nodes that reach **s** optimally are explored before nodes that reach **s** suboptimally
 - Result: A* graph search is optimal



Optimality

- Tree search:
 - A* is optimal if heuristic is admissible
 - UCS is a special case ($h = 0$)
- Graph search:
 - A* optimal if heuristic is consistent
 - UCS optimal ($h = 0$ is consistent)
- Consistency implies admissibility
- In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems


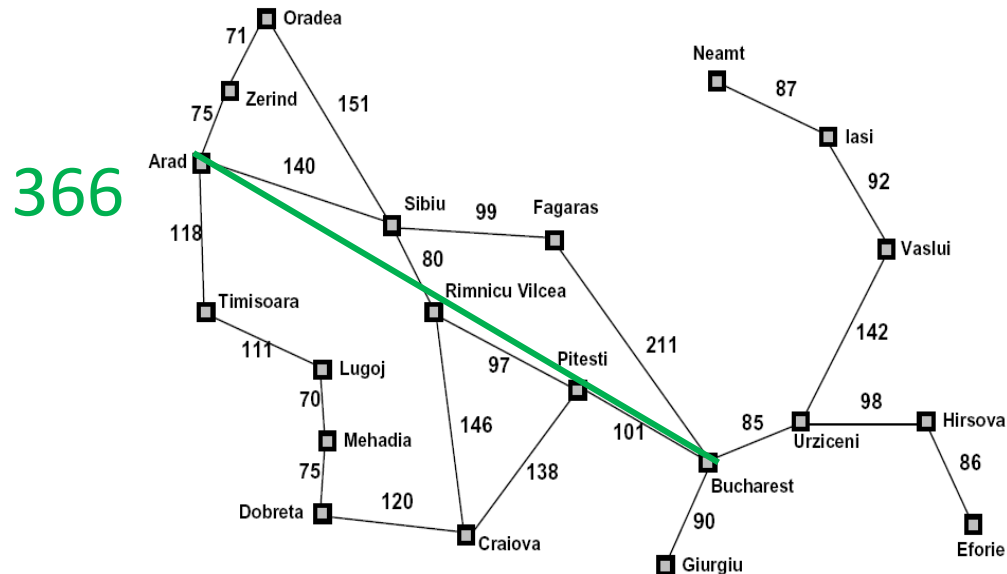


Creating Heuristics



Creating Admissible Heuristics


- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available



Example: 8 Puzzle

7	2	4
5		6
8	3	1

Start State



Actions

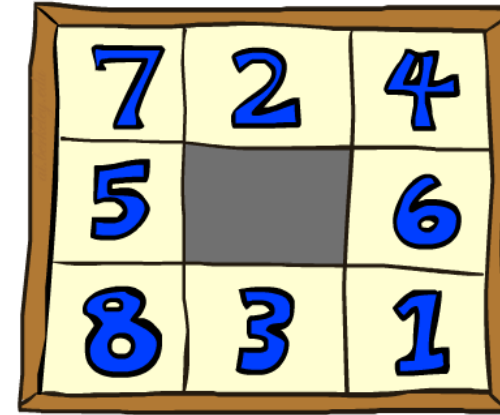
	1	2
3	4	5
6	7	8

Goal State


- What are the states?
- How many states?
- What are the actions?
- How many actions from the start state?
- What should the step costs be?

8 Puzzle I


- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) = 8$
- This is a *relaxed-problem* heuristic



Start State



Goal State



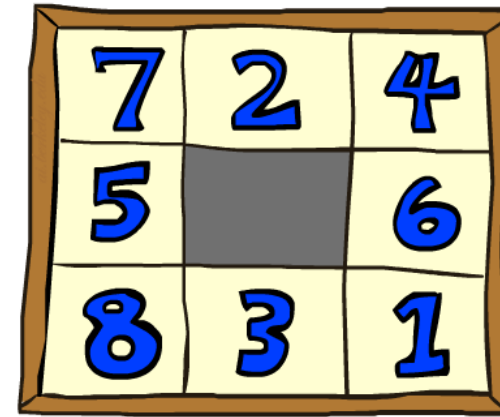
Average nodes expanded when the optimal path has...

	...4 steps	...8 steps	...12 steps
UCS	112	6,300	3.6×10^6
A*TILES	13	39	227

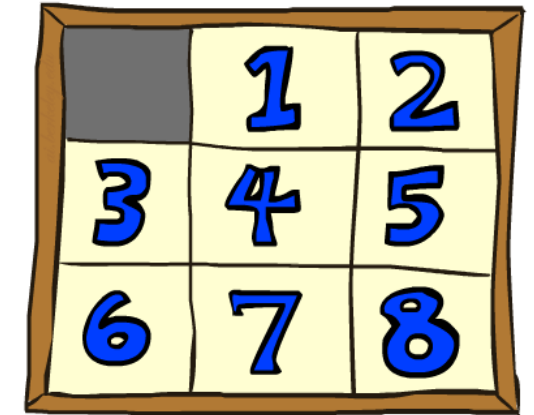
Statistics from Andrew Moore

8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?
- Total *Manhattan* distance
- Why is it admissible?
- $h(\text{start}) = 3 + 1 + 2 + \dots = 18$



Start State



Goal State

Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
A*TILES	13	39	227
A*MANHATTAN	12	25	73

Combining heuristics

- Dominance: $h_a \geq h_c$ if

$$\forall n \quad h_a(n) \geq h_c(n)$$


- Roughly speaking, larger is better as long as both are admissible
 - The **zero heuristic** is pretty bad (what does A* do with $h=0$?)
 - The **exact heuristic** is pretty good, but usually too expensive!
-
- What if we have two heuristics, neither dominates the other?
 - Form a new heuristic by taking the max of both:
$$h(n) = \max(h_a(n), h_b(n))$$
 - Max of admissible heuristics is admissible and dominates both!

A*: Summary



A*: Summary

- A* uses both backward costs and (estimates of) forward costs
- A* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems



In-Class Activity

- Q1: Practice creating heuristics and running Greedy and A* search
- Q2: Walk through Amazon Robot Example



- Notes added on LMS
- <https://www.oreilly.com/library/view/graph-algorithms/9781492047674/ch04.html>

