


Dr. Seemab latif
Lecture 5
9th Oct 2023

AI: Representation and Problem Solving

Adversarial Search



Slide credits: Pat Virtue, <http://ai.berkeley.edu>

Outline

- History / Overview
- Zero-Sum Games (Minimax)
- Evaluation Functions
- Search Efficiency (α - β Pruning)
- Games of Chance (Expectimax)



Checkers:

- 1950: First computer player.
- 1959: Samuel's self-taught program.
- 1994: First computer world champion: Chinook ended 40-year-reign of human champion Marion Tinsley using complete 8-piece endgame.
- 2007: Checkers solved! Endgame database of 39 trillion states


Chess:

- 1945-1960: Zuse, Wiener, Shannon, Turing, Newell & Simon, McCarthy.
- 1960s onward: gradual improvement under "standard model"
- 1997: special-purpose chess machine Deep Blue defeats human champion Gary Kasparov in a six-game match. Deep Blue examined 200M positions per second and extended some lines of search up to 40 ply. Current programs running on a PC rate > 3200 (vs 2870 for Magnus Carlsen).

Go:


- 1968: Zobrist's program plays legal Go, barely (b>300!)
- 2005-2014: Monte Carlo tree search enables rapid advances: current programs beat strong amateurs, and professionals with a 3-4 stone handicap.
- 2015: AlphaGo from DeepMind beats Lee Sedol

Game Playing State-of-the-Art




Types of Games

- Deterministic or stochastic?
- Perfect information (fully observable)?
- One, two, or more players?
- Turn-taking or simultaneous?
- Zero sum?




Zero-Sum Games



Zero-Sum Games

- Agents have **opposite** utilities
- Pure competition:
 - One **maximizes**, the other **minimizes**




General Games


- Agents have **independent** utilities
- Cooperation, indifference, competition, shifting alliances, and more are all possible

Standard Games


- Standard games are deterministic, observable, two-player, turn-taking, zero-sum
- Game formulation:
 - Initial state: s_0
 - Players: $\text{Player}(s)$ indicates whose move it is
 - Actions: $\text{Actions}(s)$ for player on move
 - Transition model: $\text{Result}(s,a)$
 - Terminal test: $\text{Terminal-Test}(s)$
 - Terminal values: $\text{Utility}(s,p)$ for player p
 - Or just $\text{Utility}(s)$ for player making the decision at root



Adversarial Search

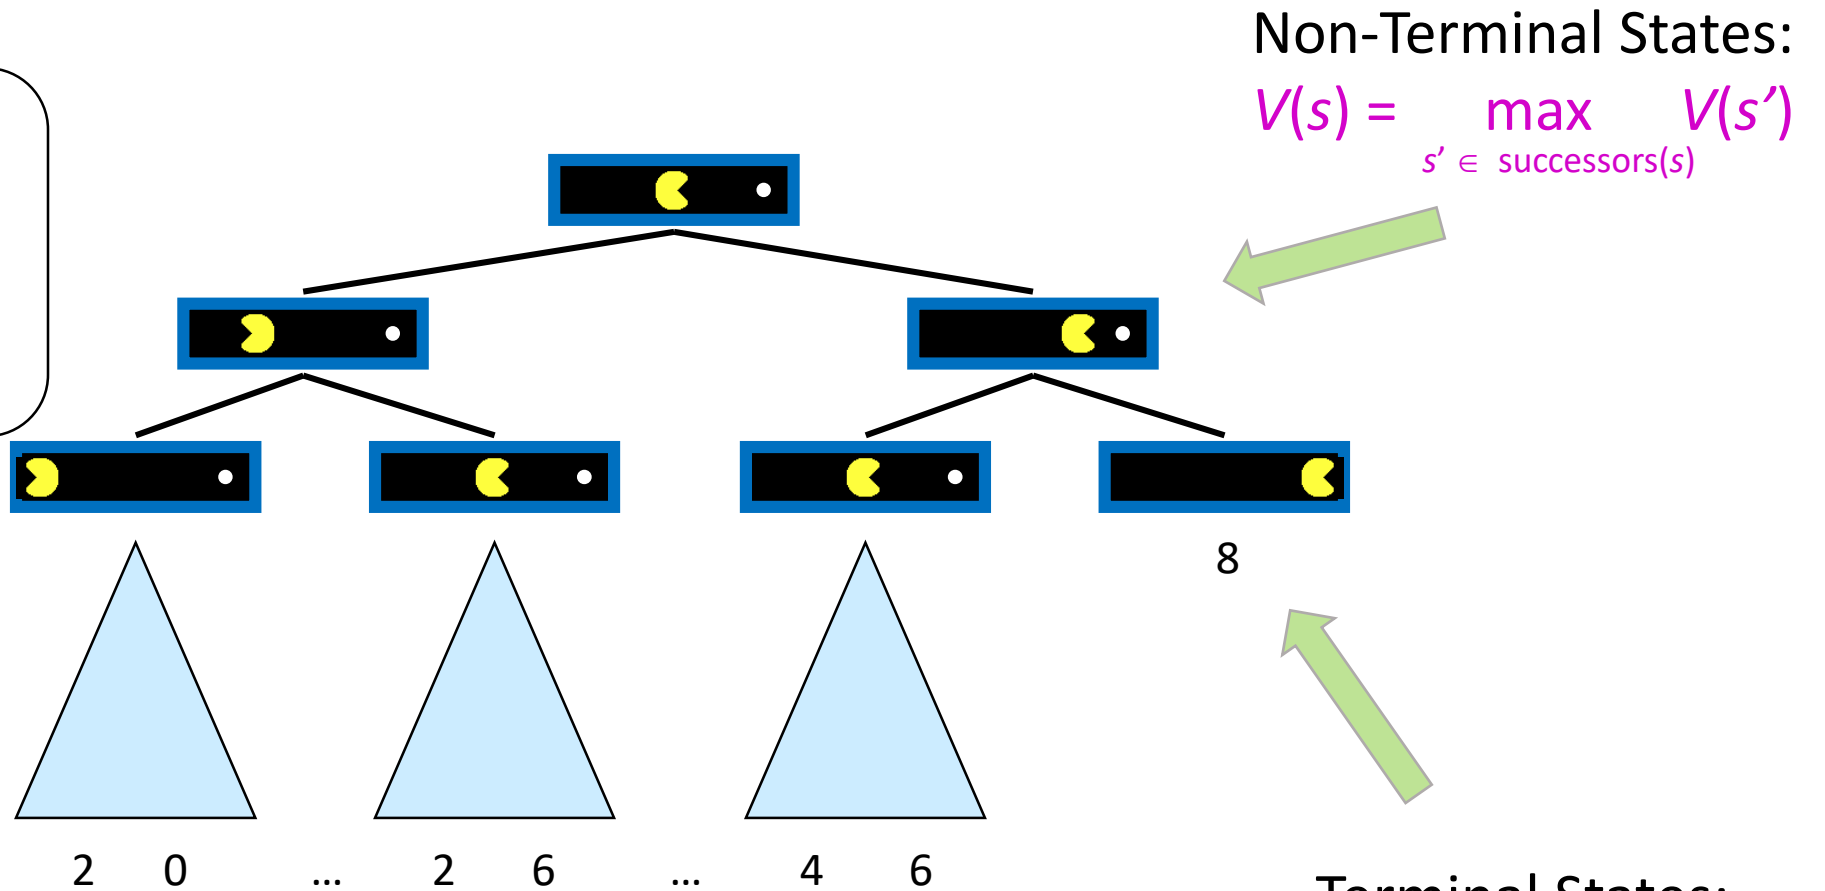


Single-Agent Trees




Value of a State

Value of a state:
The best achievable
outcome (utility)
from that state




Idea 1: Many Single-Agent Trees



Idea 2: Joint State/Action Spaces

Combine the states and actions of the N agents


$$s_0 = (s_0^A, s_0^B)$$



Idea 2: Joint State/Action Spaces

Combine the states and actions of the N agents

$$S_K = (S_K^A, S_K^B)$$




Idea 2: Joint State/Action Spaces

Search looks through all combinations of all agents' states and actions

Think of one brain controlling many agents

$$S_K = (S_K^A, S_K^B)$$



Idea 2: Joint State/Action Spaces


Search looks through all combinations of all agents' states and actions

Think of one brain controlling many agents

What is the size of the state space?

What is the size of the action space?

What is the size of the search tree?



Idea 3: Centralized Decision Making


Each agent proposes their actions and computer confirms the joint plan

Example: Autonomous driving through intersections

Idea 4: Alternate Searching One Agent at a Time

Search one agent's actions from a state, search the next agent's actions from those resulting states, etc...

Choose the best cascading combination of actions




Idea 4: Alternate Searching One Agent at a Time

Search one agent's actions from a state, search the next agent's actions from those resulting states , etc...

What is the size of the state space?


What is the size of the action space?

What is the size of the search tree?





Minimax

States
Actions
Values





Tic-Tac-Toe Game Tree



Tic-Tac-Toe Game Tree

This is a zero-sum game,
the best action for X is
the worst action for O
and vice versa

How do we define
best and worst?



Tic-Tac-Toe Game Tree



MAX (X)




MIN (O)




MAX (X)



MIN (O)



Instead of taking the max utility at every level, alternate max and min



Tic-Tac-Toe Minimax



MAX (X)




MIN (O)



MAX (X)



MIN (O)




Small Pacman Example

MAX nodes: under Agent's control

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

MIN nodes: under Opponent's control

$$V(s) = \min_{s' \in \text{successors}(s)} V(s')$$



Terminal States:

$$V(s) = \text{known}$$


Minimax Code

```
def max_value(state):  
    if state.is_leaf:  
        return state.value  
    # TODO Also handle depth limit  
  
    best_value = -10000000  
  
    for action in state.actions:  
        next_state = state.result(action)  
  
        next_value = min_value(next_state)  
  
        if next_value > best_value:  
            best_value = next_value  
  
    return best_value  
  
def min_value(state):
```




What is the minimax value at the root

1. 2
2. 3
3. 6
4. 12
5. 14





1-answer





What kind of search is Minimax Search?

- A) BFS
- B) DFS
- C) UCS
- D) A*




2-answer

What kind of search is Minimax Search?

- A) BFS
- B) DFS**
- C) UCS
- D) A*


Minimax Efficiency

- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$
- Example: For chess, $b \approx 35$, $m \approx 100$
 - Exact solution is completely infeasible
 - Humans can't do this either, so how do we play chess?



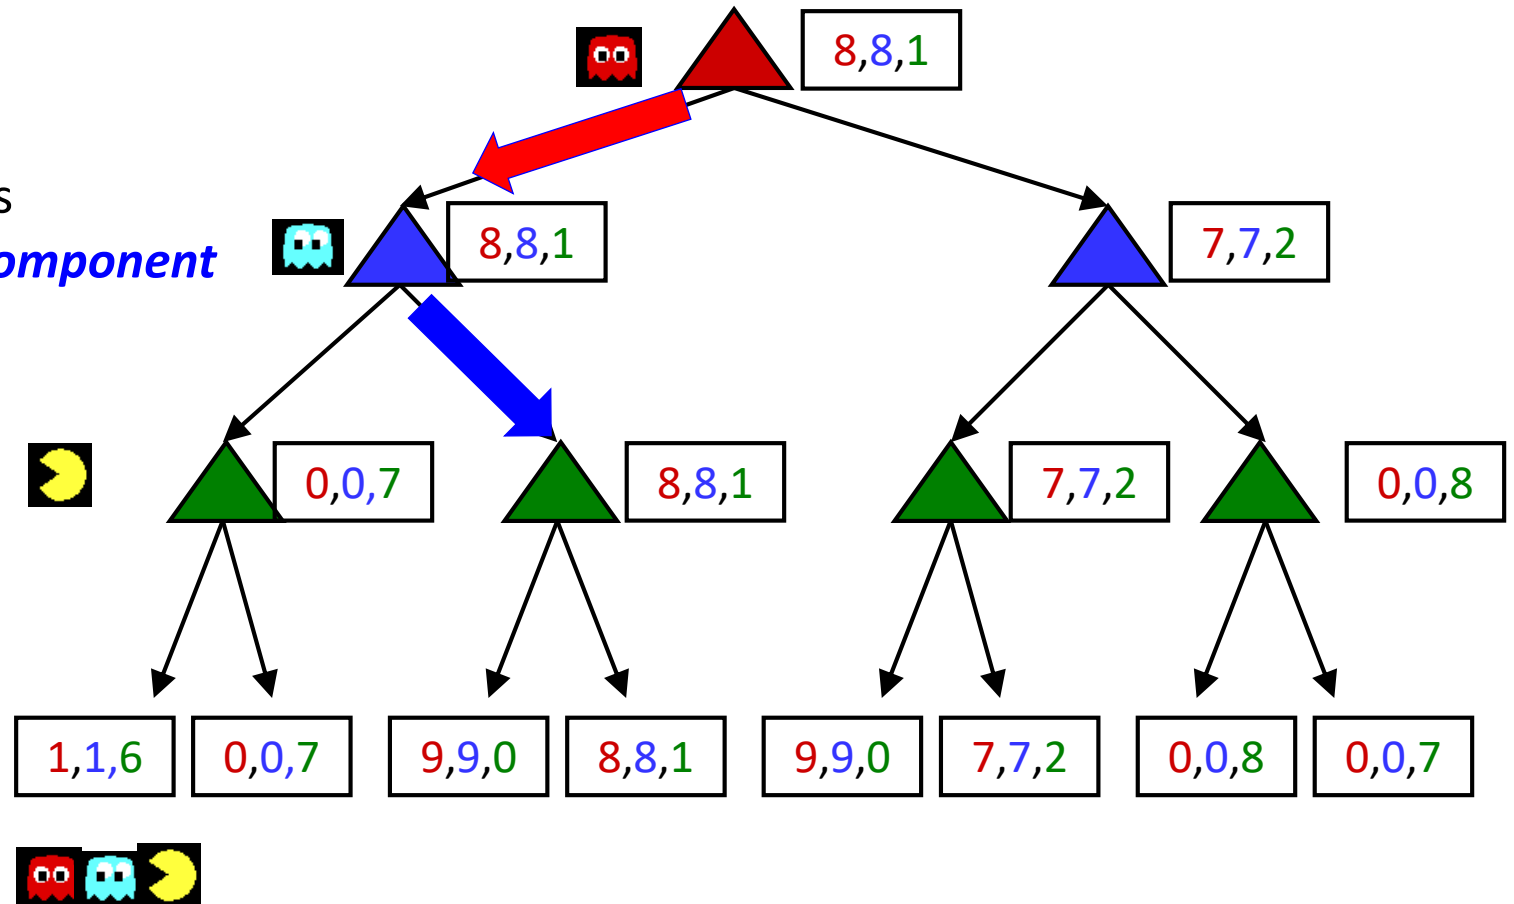

Small Size Robot Soccer

- Joint State/Action space and search for our team
- Adversarial search to predict the opponent team



Generalized minimax

- What if the game is not zero-sum, or has multiple players?
- Generalization of minimax:
 - Terminals have **utility tuples**
 - Node values are also utility tuples
 - **Each player maximizes its own component**
 - Can give rise to cooperation and competition dynamically...




Resource Limits




Resource Limits

- Problem: In realistic games, cannot search to leaves!
- Solution 1: Bounded lookahead
 - Search only to a preset **depth limit** or **horizon**
 - Use an **evaluation function** for non-terminal positions
- Guarantee of optimal play is gone
- More plies make a BIG difference
- Example:
 - Suppose we have 100 seconds, can explore 10K nodes / sec
 - So can check 1M nodes per move
 - For chess, $b \sim 35$ so reaches about depth 4 – not so good




Depth Matters

- Evaluation functions are always imperfect
- Deeper search => better play (usually)
- Or, deeper search gives same quality of play with a less accurate evaluation function
- An important example of the tradeoff between complexity of features and complexity of computation




Evaluation Functions




Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search




- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:
 - $EVAL(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$
 - E.g., $w_1 = 9$, $f_1(s) = (\text{num white queens} - \text{num black queens})$, etc.
- Terminate search only in **quiescent** positions, i.e., no major changes expected in feature values


Evaluation for Pacman



Solution 2: Game Tree Pruning




Intuition: prune the branches that can't be chosen



Alpha-Beta Pruning Example

α = best option so far from any
MAX node on this path



We can prune when: min node won't be higher than 2, while parent max has seen something larger in another branch


The order of generation matters: more pruning is possible if good moves come first

Alpha-Beta Implementation

α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$   
            return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$   
            return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```




What is the value of the blue triangle?

- A) 10
- B) 8
- C) 4
- D) 50




3-answer



What is the value of the blue triangle?

- A) 10
- B) 8**
- C) 4
- D) 50

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

if $v \geq \beta$

return v

$\alpha = \max(\alpha, v)$

return v

def min-value(state, α , β):

initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


if $v \leq \alpha$

return v

$\beta = \min(\beta, v)$

return v

Alpha-Beta Small Example



def max-value(state, α , β):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

if $v \geq \beta$

return v

$\alpha = \max(\alpha, v)$

return v

def min-value(state, α , β):

initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


if $v \leq \alpha$

return v

$\beta = \min(\beta, v)$

return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

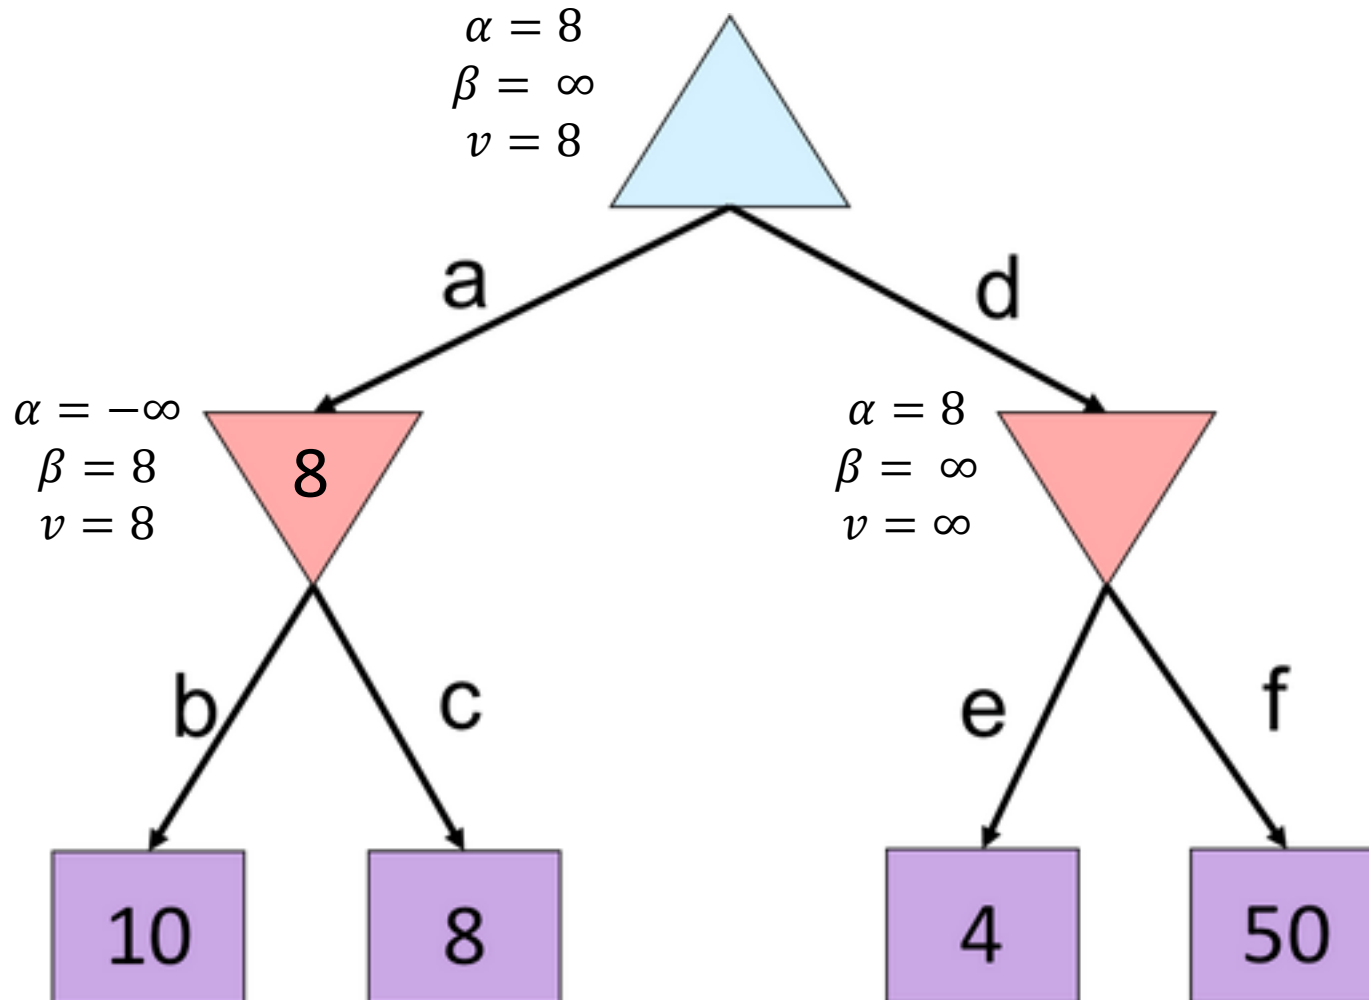
 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

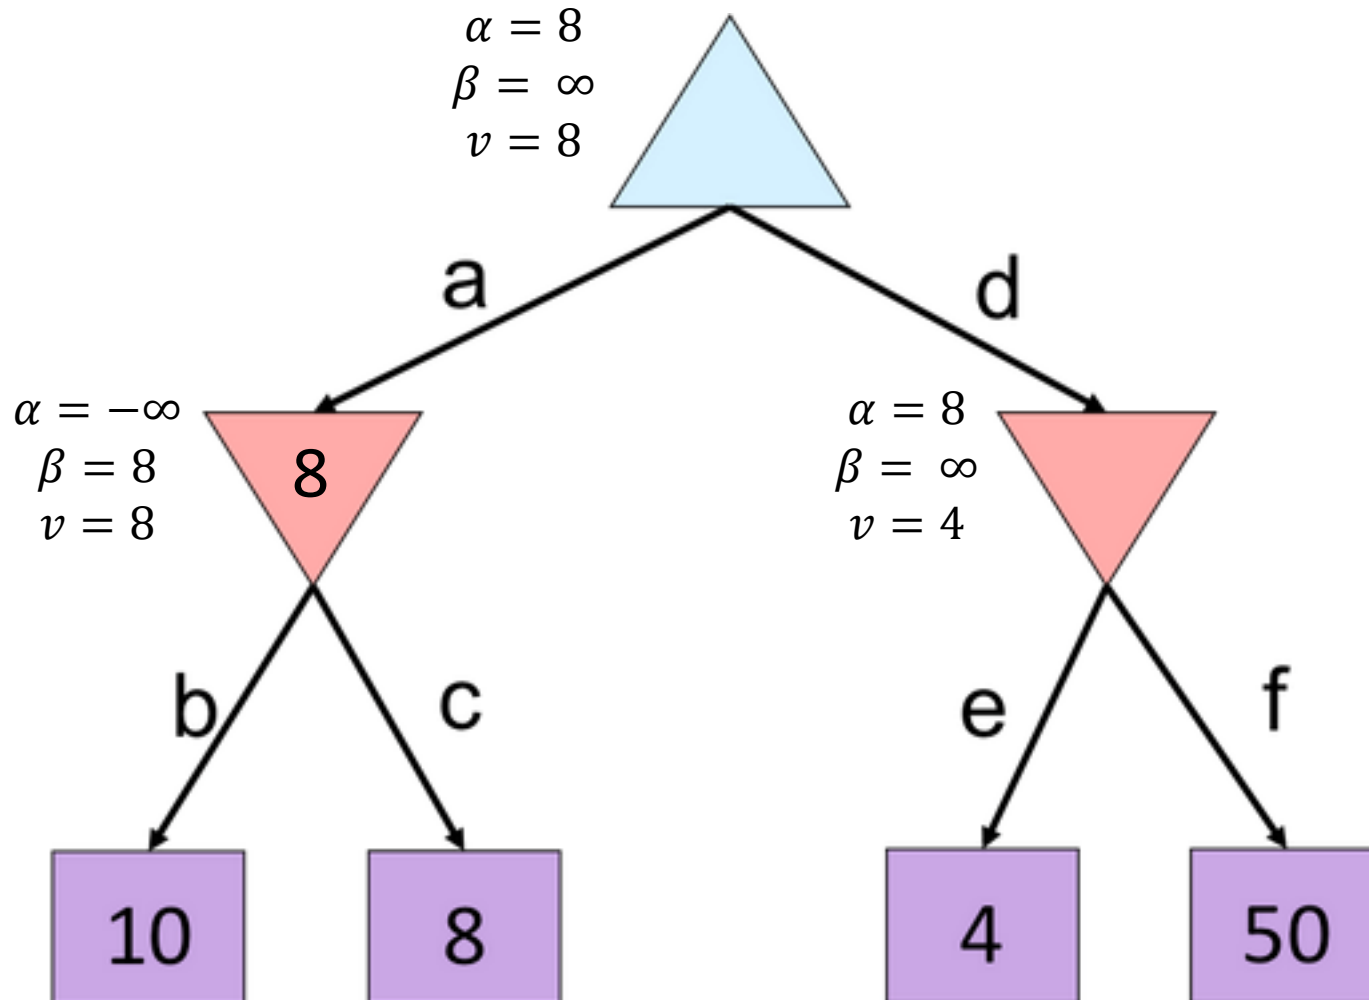
 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

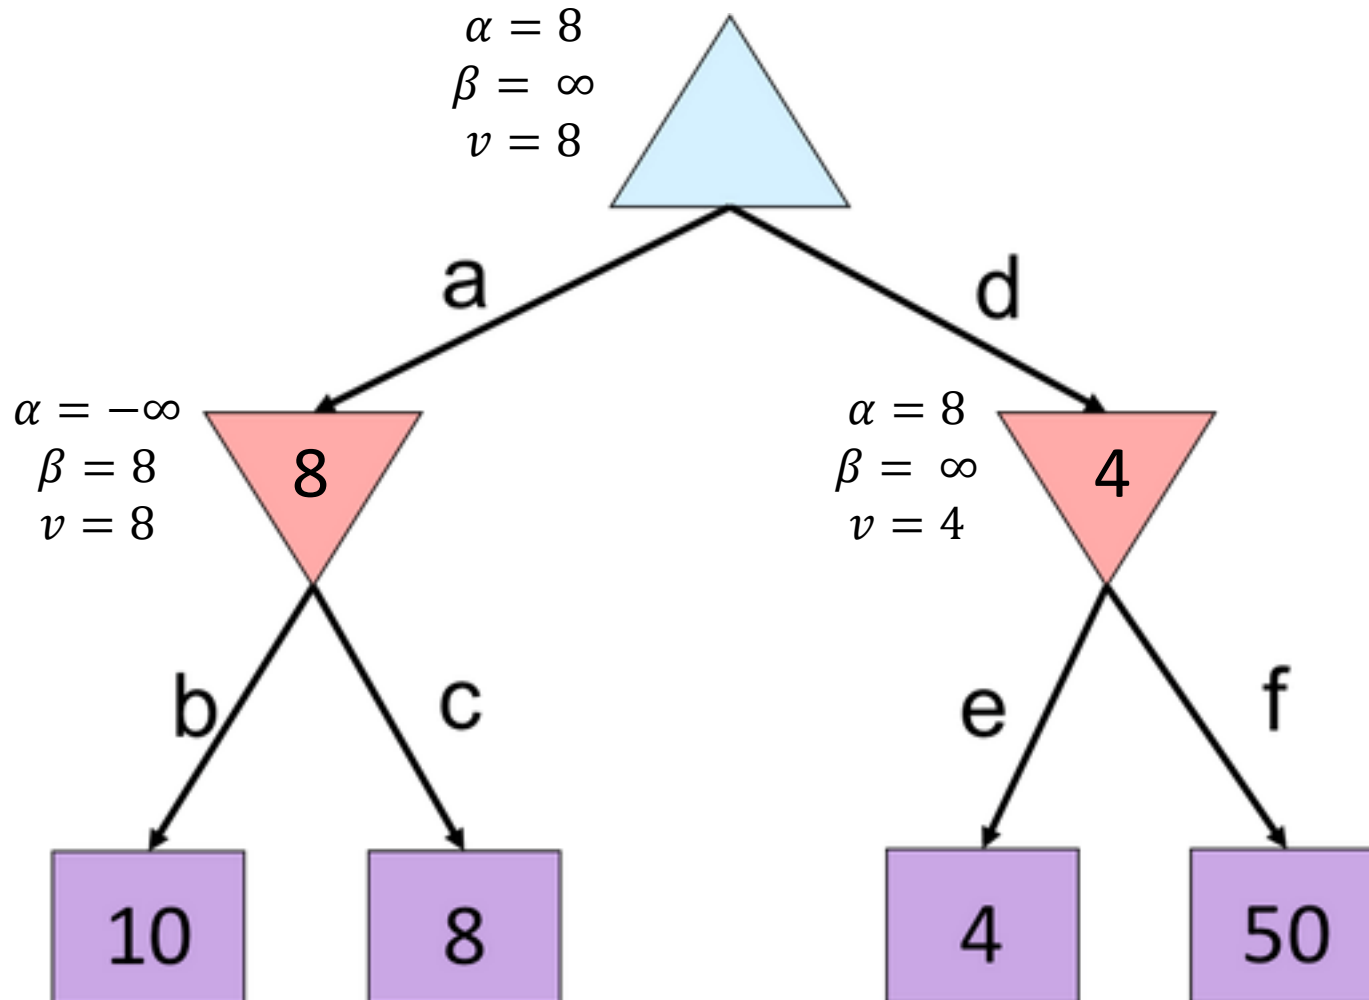
 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$


 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

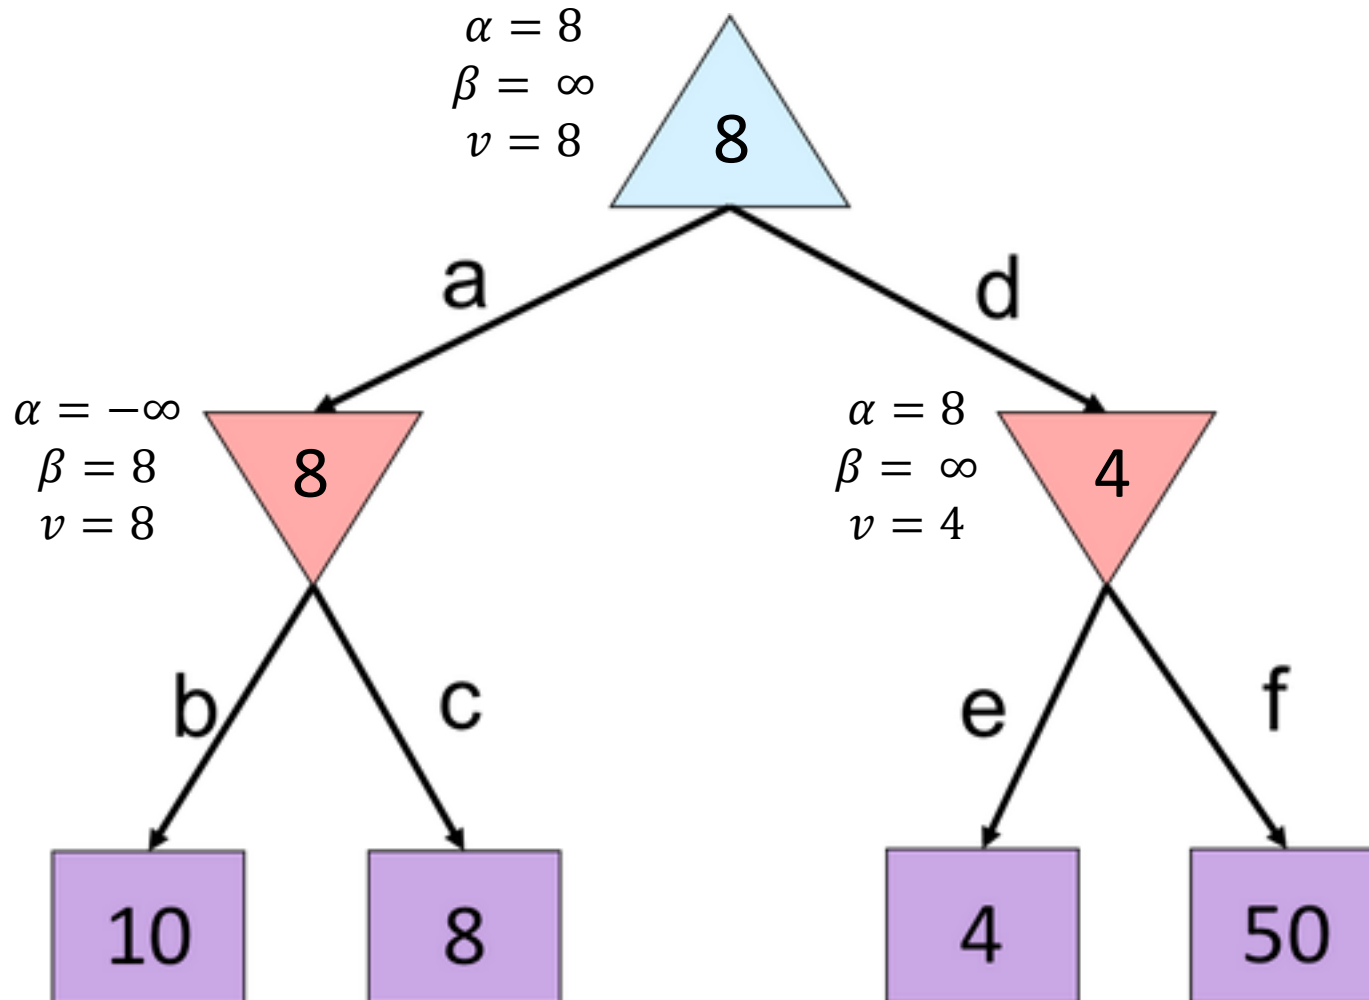
 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$

 return v

Alpha-Beta Small Example



def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$


 return v



4

What is the value of the top node?

- A) 10
- B) 100
- C) 2
- D) 4





5

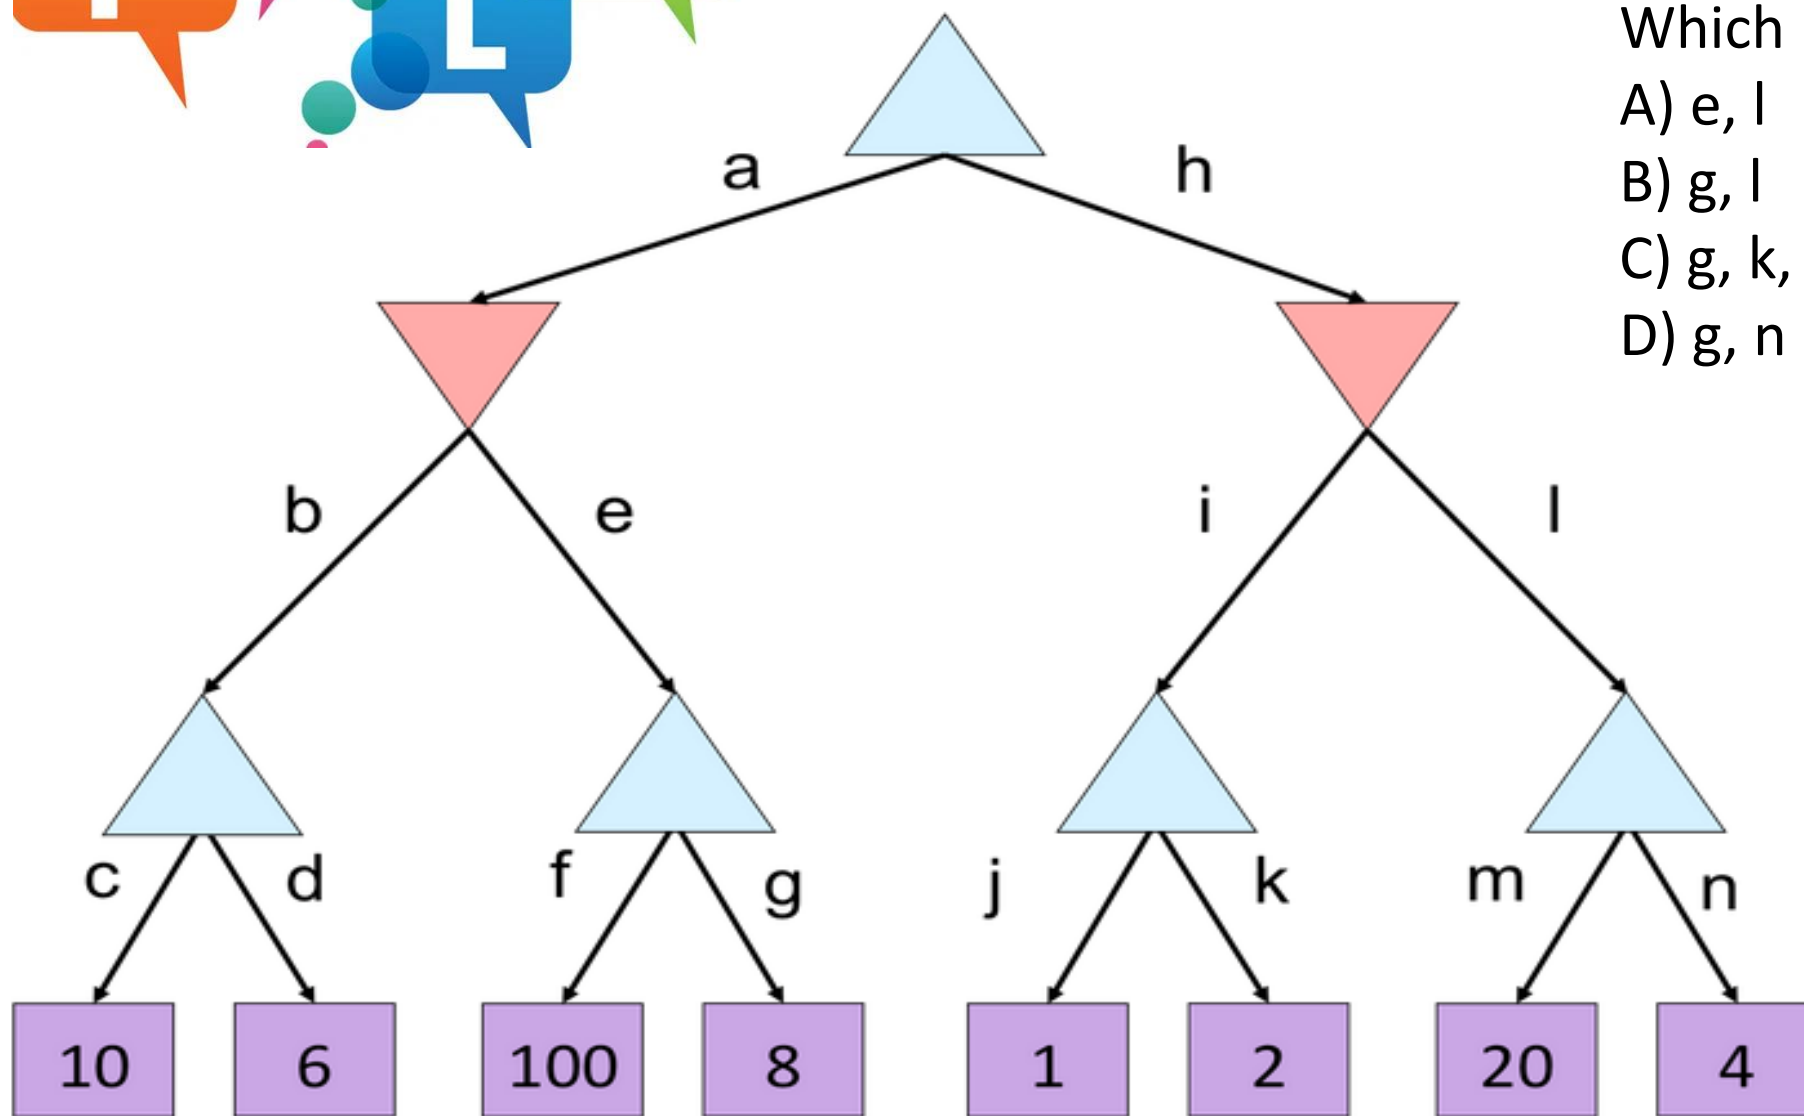
Which branches are pruned?

A) e, l

B) g, l

C) g, k, l

D) g, n





5-answer
 $\alpha = 10$


Which branches are pruned?

A) e, l

B) g, l


C) g, k, l

D) g, n




Alpha-Beta Pruning Properties

- Theorem: This pruning has **no effect** on minimax value computed for the root!
- Good child ordering improves effectiveness of pruning
 - Iterative deepening helps with this
- With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!
 - 1M nodes/move => depth=8, respectable





- This is a simple example of **metareasoning** (computing about what to compute)

Games with uncertain outcomes





Modeling Assumptions

Know your opponent



Modeling Assumptions

Know your opponent



Modeling Assumptions

Dangerous Pessimism

Assuming the worst case when it's not likely





Dangerous Optimism

Assuming chance when the world is adversarial





Chance outcomes in trees





Tictactoe, chess

Minimax



Tetris, investing

Expectimax



Backgammon, Monopoly

Expectiminimax

Minimax

function **decision(s)** returns an action

return the action **a** in **Actions(s)** with the highest
value(Result(s,a))



function **value(s)** returns a value

if **Terminal-Test(s)** then return **Utility(s)**

if **Player(s) = MAX** then return **max**_{a in Actions(s)} **value(Result(s,a))**

if **Player(s) = MIN** then return **min**_{a in Actions(s)} **value(Result(s,a))**

Expectiminimax

function **decision(s)** returns an action

return the action **a** in **Actions(s)** with the highest
value(Result(s,a))



function **value(s)** returns a value


if **Terminal-Test(s)** then return **Utility(s)**

if **Player(s) = MAX** then return **max**_{a in Actions(s)} **value(Result(s,a))**

if **Player(s) = MIN** then return **min**_{a in Actions(s)} **value(Result(s,a))**

if **Player(s) = CHANCE** then return **sum**_{a in Actions(s)} **Pr(a) * value(Result(s,a))**

Probabilities



Probabilities

A **random variable** represents an event whose outcome is unknown

A **probability distribution** is an assignment of weights to outcomes

Example: Traffic on freeway


□ Random variable: T = whether there's traffic

□ Outcomes: T in {none, light, heavy}

□ Distribution:

$$P(T=\text{none}) = 0.25, \quad P(T=\text{light}) = 0.50, \quad P(T=\text{heavy}) = 0.25$$

Probabilities over all possible outcomes sum to one



0.25



0.50






0.25

Expected Value

- The expected value of a random variable is the average, weighted by the probability distribution over outcomes
- Example: How long to get to the airport?

$$\begin{array}{rcccl} \text{Time:} & 20 \text{ min} & & 30 \text{ min} & & 60 \text{ min} \\ & \times & + & \times & + & \times \\ \text{Probability:} & 0.25 & & 0.50 & & 0.25 \end{array}$$



→ 35 min



Expectations

Time: 20 min x 0.25 + 30 min x 0.50 + 60 min x 0.25

Probability: 0.25 0.50 0.25



Max node notation

$$V(s) = \max_a V(s'),$$


where $s' \in \text{re } \text{lt}(s, a)$

Chance node notation

$$V(s) = \sum_{s'} P(s') V(s')$$


Expectimax Pseudocode

$\text{sum}_{a \text{ in Action}(s)} \text{Pr}(a) * \text{value}(\text{Result}(s,a))$




$$v = (1/2) (8) + (1/3) (24) + (1/6) (-12) = 10$$

Expectimax Example



What Values to Use?



$$x > y \Rightarrow f(x) > f(y)$$

$$f(x) = Ax + B \text{ where } A > 0$$

- For worst-case minimax reasoning, evaluation function scale doesn't matter
 - We just want better states to have higher evaluations (get the ordering right)
 - Minimax decisions are ***invariant with respect to monotonic transformations on values***
- Expectiminimax decisions are ***invariant with respect to positive affine transformations***
- Expectiminimax evaluation functions have to be aligned with actual win probabilities!

Summary

- Multi-agent problems can require more space or deeper trees to search
- Games require decisions when optimality is impossible
 - Bounded-depth search and approximate evaluation functions
- Games force efficient use of computation
 - Alpha-beta pruning
- Game playing has produced important research ideas
 - Reinforcement learning (checkers)
 - Iterative deepening (chess)
 - Rational metareasoning (Othello)
 - Monte Carlo tree search (Go)
 - Solution methods for partial-information games in economics (poker)
- Video games present much greater challenges – lots to do!
 - $b = 10^{500}$, $|S| = 10^{4000}$, $m = 10,000$